

In [32]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import rcParams
from matplotlib.cm import rainbow
%matplotlib inline
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn import datasets
from sklearn import linear_model
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeClassifier
```

In [33]:

```
data= pd.read_csv(r'C:\Users\ASUS\Downloads\archive (2)\heart.csv')
print(data)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
\										
0	52	1	0	125	212	0	1	168	0	1.0
1	53	1	0	140	203	1	0	155	1	3.1
2	70	1	0	145	174	0	1	125	1	2.6
3	61	1	0	148	203	0	1	161	0	0.0
4	62	0	0	138	294	1	1	106	0	1.9
...	...	...	..	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0
1021	60	1	0	125	258	0	0	141	1	2.8
1022	47	1	0	110	275	0	0	118	1	1.0
1023	50	0	0	110	254	0	0	159	0	0.0
1024	54	1	0	120	188	0	1	113	0	1.4

	slope	ca	thal	target
0	2	2	3	0
1	0	0	3	0
2	0	0	3	0
3	2	1	3	0
4	1	3	2	0
...	...	..	...	...
1020	2	0	2	1
1021	1	1	3	0
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

[1025 rows x 14 columns]

In [34]:

```
data.head()
```

Out[34]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	

In [35]:

```
data.tail()
```

Out[35]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3

In [36]:

```
data.shape
```

Out[36]:

(1025, 14)

In [37]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

In [38]:

```
data.isnull().sum() #to check for null values
```

Out[38]:

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

In [39]:

```
data.describe()
```

Out[39]:

	age	sex	cp	trestbps	chol	fbs	re
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.500000
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.500000
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000




In [40]:

```
plt.figure(figsize = (10,10))
sns.heatmap(data.corr(),annot = True,cmap='viridis',linewidths=.5)
```

Out[40]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x2caeb9d3760>




In [41]:

```
data.hist()
```

Out[41]:

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB11FEE
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB819A3
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB93A7C
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAECA0776
0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB9B46A
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB9EF07
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB9EF25
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAEB8AC3A
0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000002CAECFB99D
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED078E2
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED0B22B
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED0DC70
0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED109B8
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED137FD
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED1703A
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000002CAED1909A
0>]],
      dtype=object)
```



In [42]:

```
data['target'].value_counts() #how many ppl have heart disease
```

Out[42]:


```
1    526
0    499
Name: target, dtype: int64
```

In [43]:

```
rcParams['figure.figsize'] = 8,6
plt.bar(data['target'].unique(),data['target'].value_counts(),color=['red','green'])
plt.xticks([0,1])
plt.xlabel('Target Classes')
plt.ylabel('Count')
plt.title('Count of each Target Class')
```

Out[43]:

Text(0.5, 1.0, 'Count of each Target Class')



In [44]:

```
x = data.drop(columns = 'target',axis = 1) #dropping row axis = 0
y = data['target']
```

In [45]:

```
print(x)
```

	age	sex	cp	trestbps	chol	fb	restecg	thalach	exang	oldpeak
\										
0	52	1	0	125	212	0	1	168	0	1.0
1	53	1	0	140	203	1	0	155	1	3.1
2	70	1	0	145	174	0	1	125	1	2.6
3	61	1	0	148	203	0	1	161	0	0.0
4	62	0	0	138	294	1	1	106	0	1.9
...	...	...	..	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0
1021	60	1	0	125	258	0	0	141	1	2.8
1022	47	1	0	110	275	0	0	118	1	1.0
1023	50	0	0	110	254	0	0	159	0	0.0
1024	54	1	0	120	188	0	1	113	0	1.4

	slope	ca	thal
0	2	2	3
1	0	0	3
2	0	0	3
3	2	1	3
4	1	3	2
...	...	..	...
1020	2	0	2
1021	1	1	3
1022	1	1	2
1023	2	0	2
1024	1	1	3

[1025 rows x 13 columns]

In [46]:

```
print(y)
```

0	0
1	0
2	0
3	0
4	0
...	..
1020	1
1021	0
1022	0
1023	1
1024	0

Name: target, Length: 1025, dtype: int64

In [47]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,stratify=y,random_state = 2)
```



In [49]:

```
print(x.shape,x_train.shape,x_test.shape)
```

(1025, 13) (820, 13) (205, 13)

In [72]:

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(x_train,y_train)
```

Out[72]:

LinearRegression()

In [73]:

```
y_predicted = model.predict(x_test)
```

In [74]:

```
print("MSE is:=",mean_squared_error(y_test,y_predicted))
```

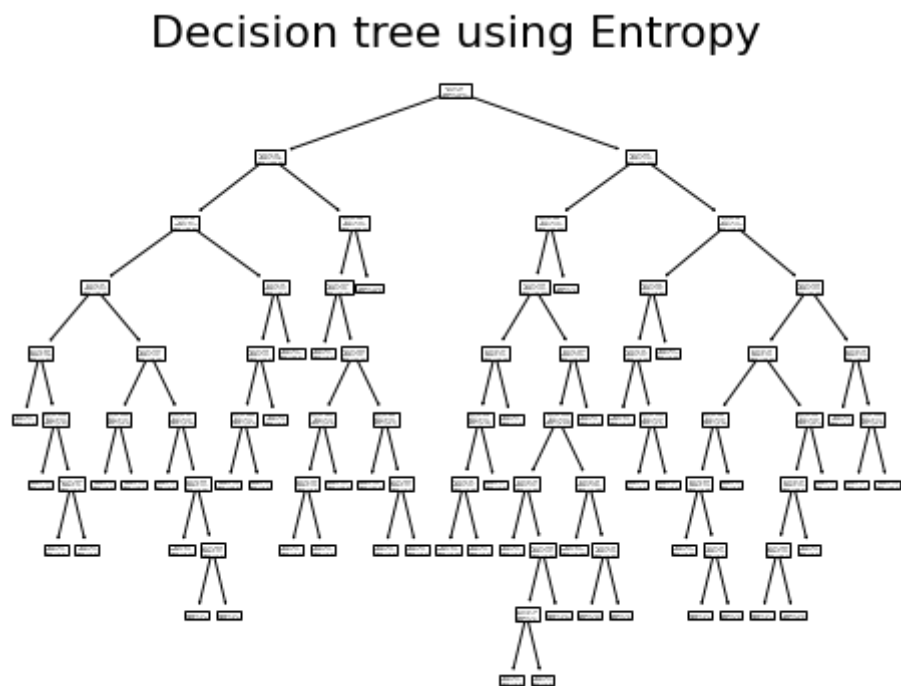
MSE is:= 0.13612944595506918

In [75]:

```
model1 = DecisionTreeClassifier()
model1.fit(x_train,y_train)
prediction = model1.predict(x_test)
import sklearn as sklearn
sklearn.tree.plot_tree(model1,max_depth=None,feature_names = None,class_names=None,label='all',filled=False,impurity=True,node_ids=False,proportion=False,rounded=False,precision=3,ax=None,fontsize=None)
plt.title('Decision tree using Entropy',fontsize=22)
```

Out[75]:

```
Text(0.5, 1.0, 'Decision tree using Entropy')
```



In [76]:

```
score = accuracy_score(y_test,prediction)
score
```

Out[76]:

1.0

In [77]:

```
model2 = LogisticRegression()
```

In [78]:

```
model2.fit(x_train,y_train) #find relationship btw feature and target
```

C:\anaconda\lib\site-packages\sklearn\linear\_model\\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Out[78]:

LogisticRegression()

In [70]:

```
#Accuracy on training data
x_train_prediction = model2.predict(x_train)
training_data_accuracy = accuracy_score(x_train_prediction,y_train)
print('accuracy on training data:',training_data_accuracy)
```

accuracy on training data: 0.848780487804878

In [59]:

```
#Accuracy on testing data
x_test_prediction = model2.predict(x_test)
testing_data_accuracy = accuracy_score(x_test_prediction,y_test)
print('accuracy on training data:',testing_data_accuracy)
```

accuracy on training data: 0.8048780487804879

In [60]:

```
#predictive system

input_data = (71,0,0,112,149,0,1,125,0,1.6,1,0,2)
#change the ip data to numpy array
input_data_as_numpy_array = np.asarray(input_data)

#reshape the numpy array as we are predicting for only one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model2.predict(input_data_reshaped)
print(prediction)

if(prediction[0]==0):
    print("The person does not has heart disease.")
else:
    print("The person has heart disease.")
```

```
[1]
The person has heart disease.
```